



Welcome Everybody



Almutaz Adileh
Cansu Kaynak
Pejman Lotfi-Kamran
Stavros Volos








CloudSuite on Flexus


- CloudSuite: Suite for scale-out datacenter services
- Flexus: Fast, accurate & flexible architectural Simulator
- Now, CloudSuite Simics images (up to 64 cores)
- The tutorial is interactive
 - Please ask questions anytime during tutorial



2



Agenda



-  CloudSuite benchmarks overview
-  Full-system simulation with Simics
-  Flexus internals and hands-on
-  Fast simulation via statistical sampling
-  Using CloudSuite images






CloudSuite: A Suite for Emerging Scale-out Applications

Almutaz Adileh





Clouds are Scale-out





- Cloud computing is pervasive
 - User base growing exponentially
 - New services appearing daily
- Serving a global-scale audience requires scaling-out
 - Distribute data and computation to many servers
- How to characterize popular scale-out applications?



Need scale-out benchmarks

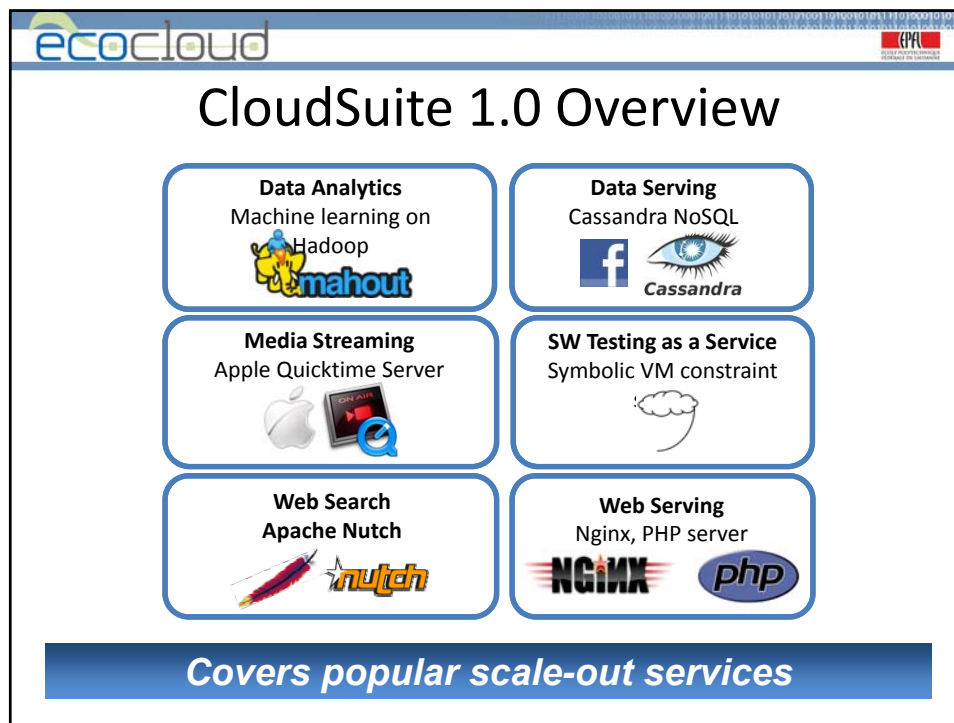
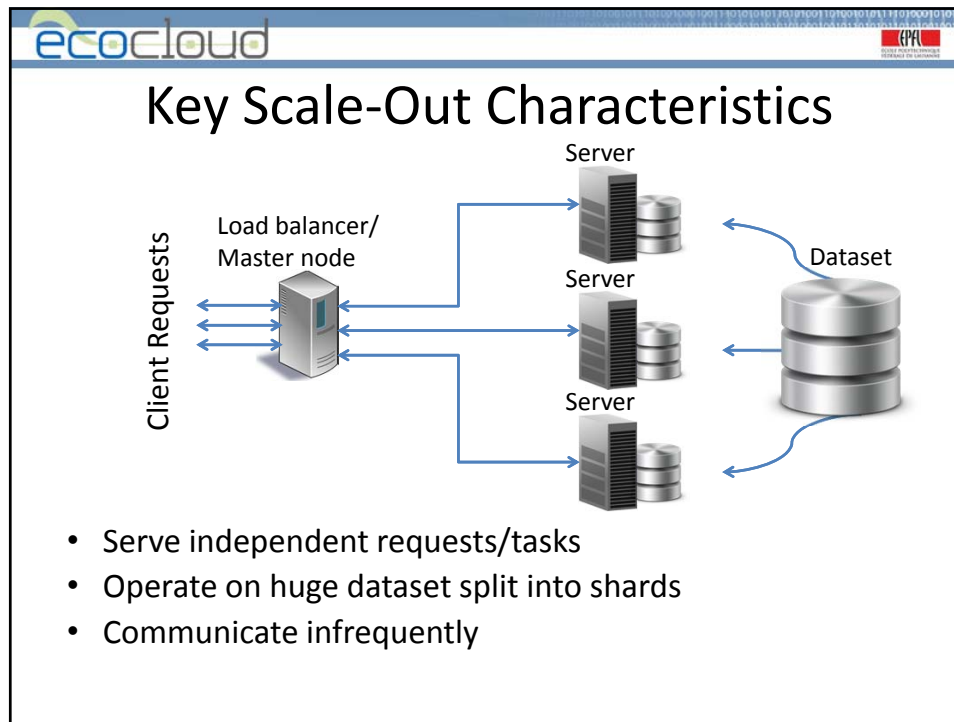




Which Benchmarks to Use?



- Benchmarks designed for scale-up


Don't represent scale-out applications







Data Analytics


- Service fast data generation rates (Big Data)
- Extract useful information from data
 - Predict user preferences, opinions, behavior
 - Benefit from information (e.g., business, security)
- Several examples
 - Book recommendation (Amazon)
 - Spyware detection (Facebook)
 - Photo interestingness (Flickr)

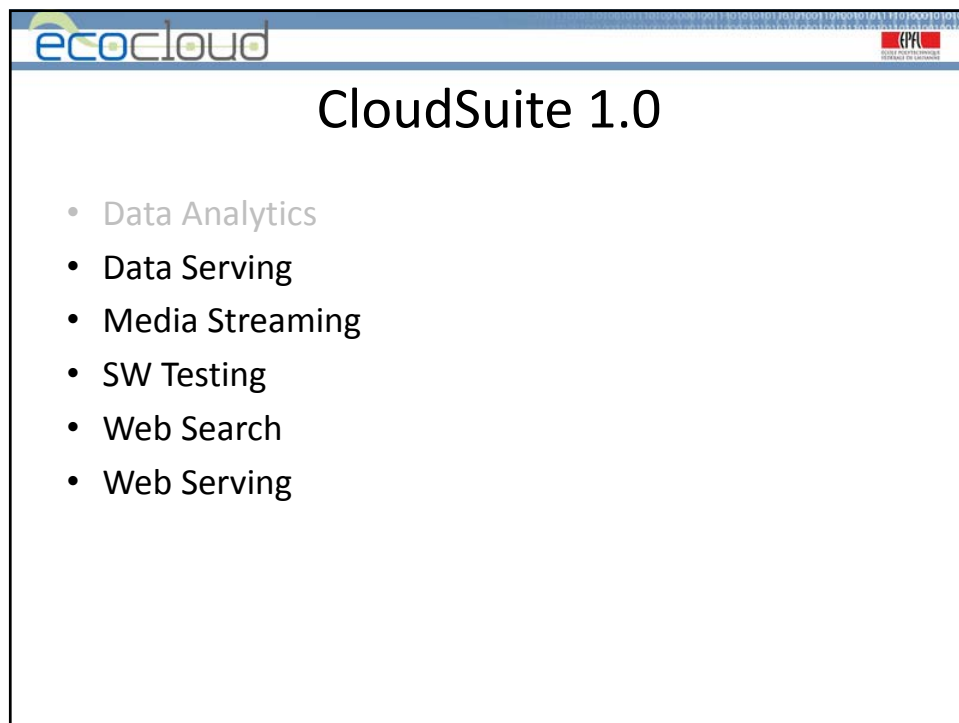
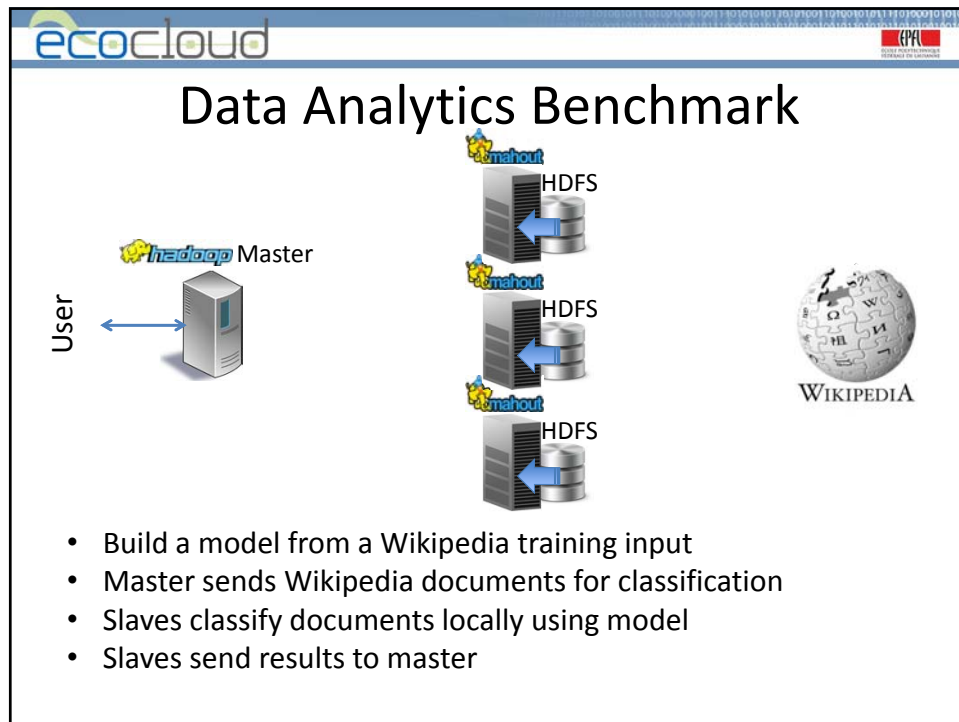


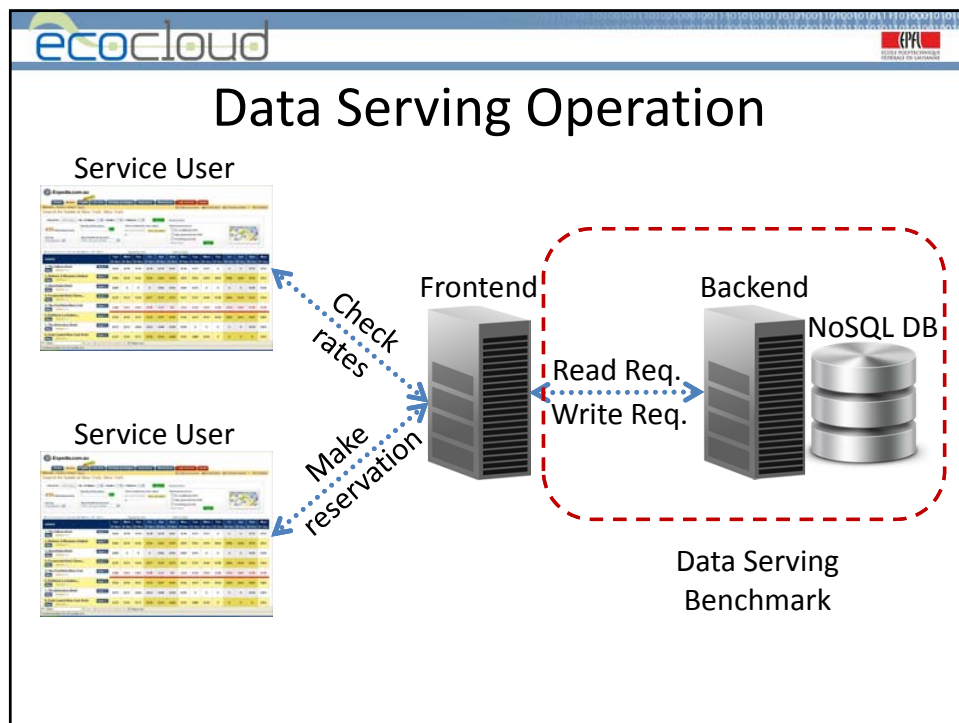
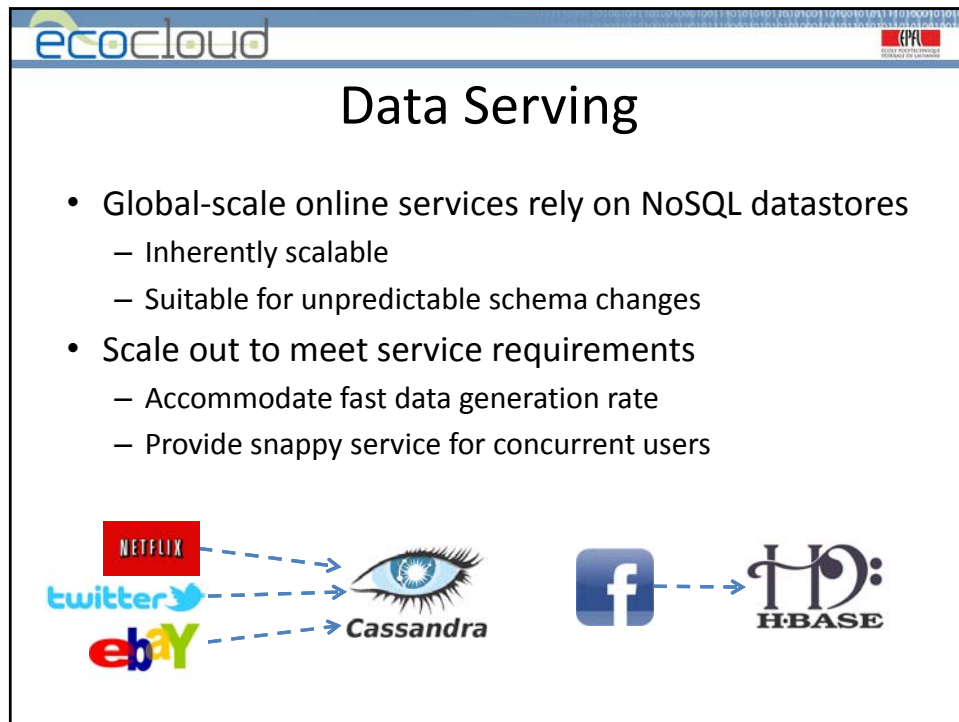


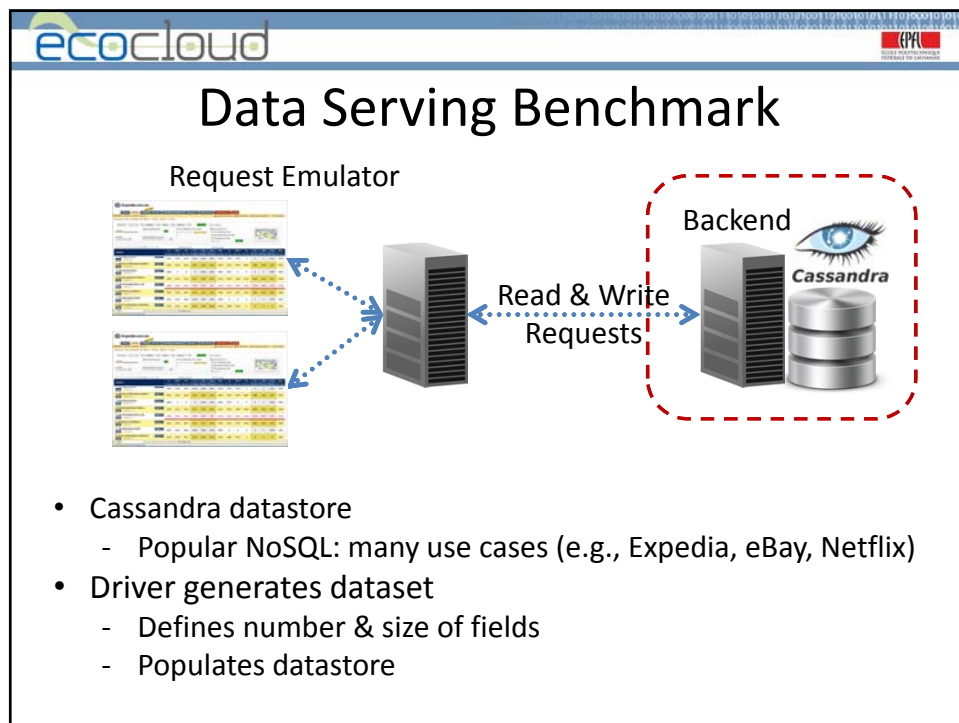
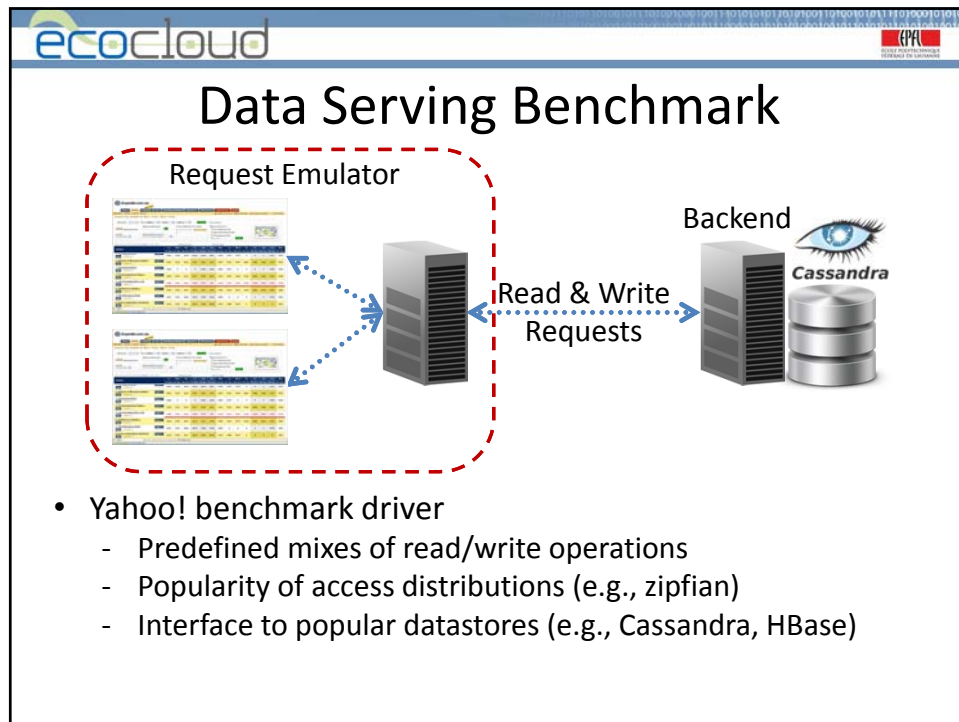
Data Analytics Benchmark



- **Application:** Text classification
 - Sentiment Analysis
 - Spam Identification
- **Software:** Mahout (Apache)
 - Popular MapReduce machine learning library
- **Dataset:** Wikipedia English page articles











CloudSuite 1.0

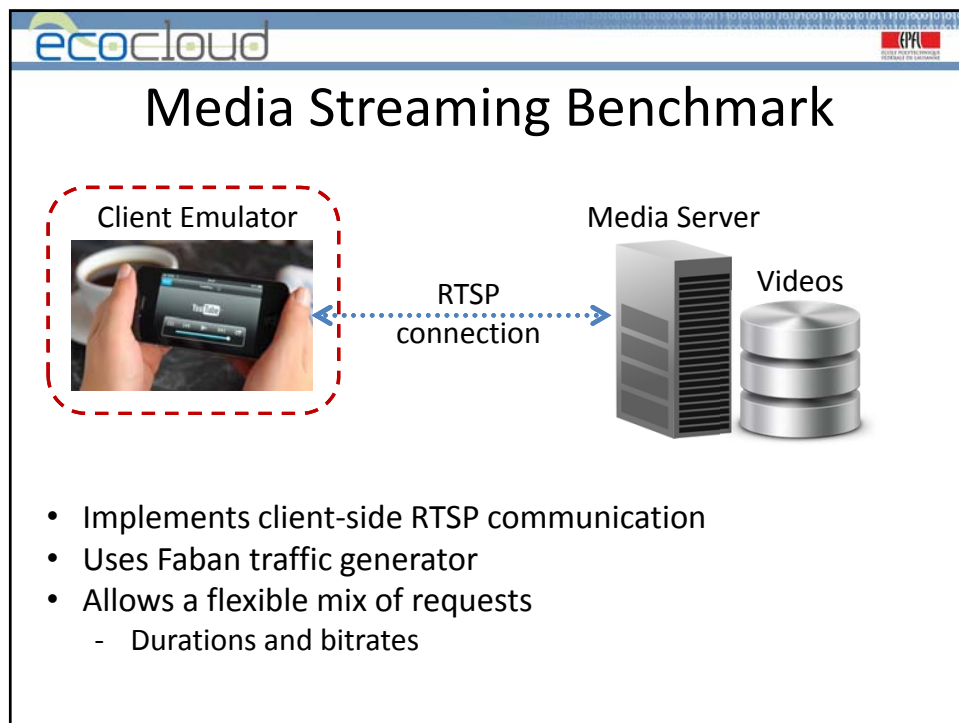
- Data Analytics
- Data Serving
- Media Streaming
- SW Testing
- Web Search
- Web Serving

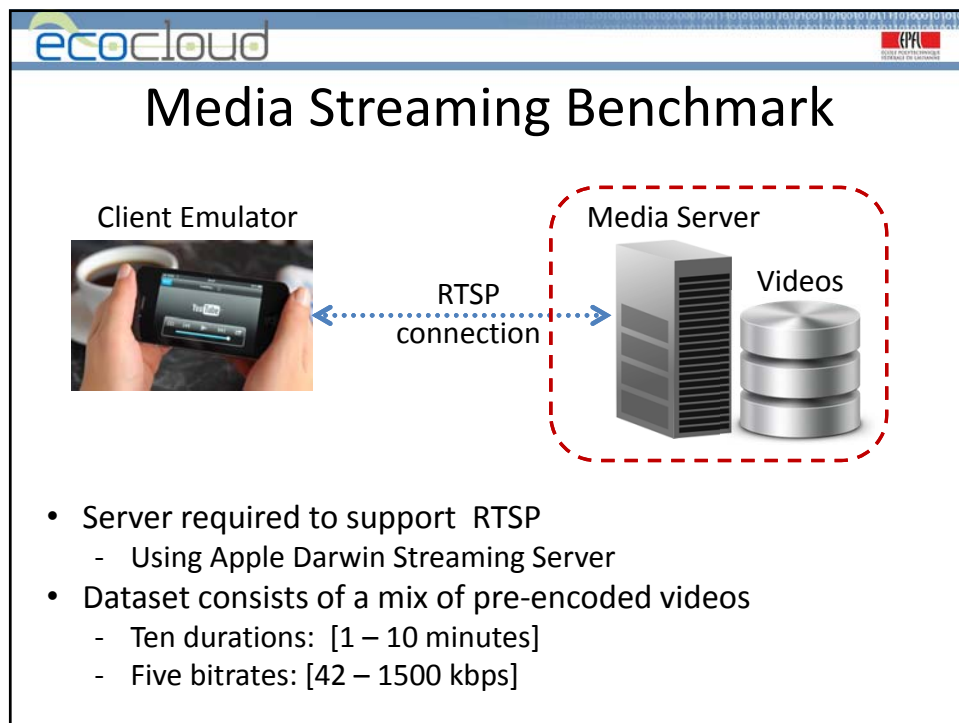
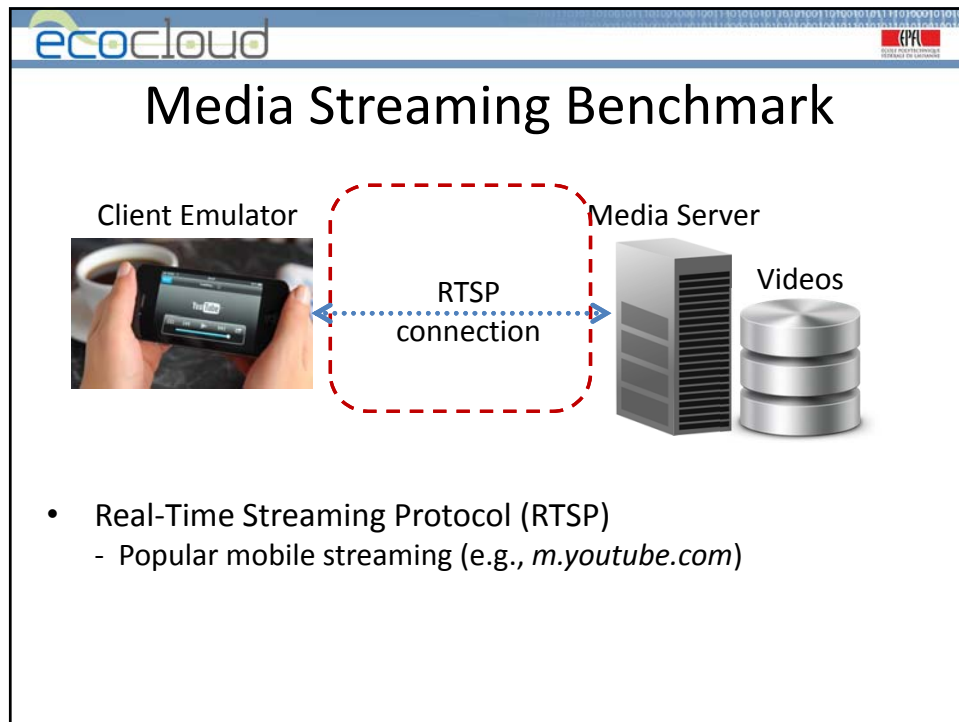




Media Streaming

- Media streaming expected to dominate internet traffic
- Increasing popularity of media streaming services
 - Video sharing sites, movie streaming services, etc.













CloudSuite 1.0

- Data Analytics
- Data Serving
- Media Streaming
- SW Testing
- Web Search
- Web Serving

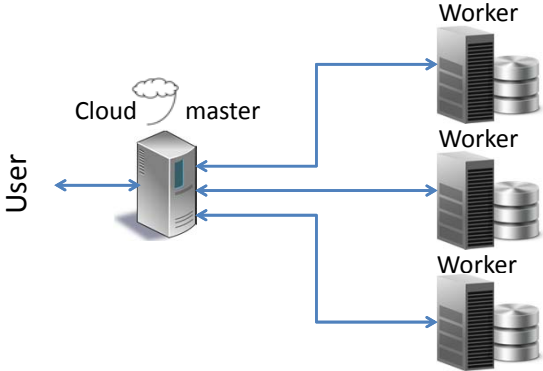


Software Testing

- Clouds allow dynamic resource allocation as needed
 - Enables previously impossible engineering practices
- Software Testing as a Service leverages cloud env.
 - Large-scale symbolic execution for SW testing
 - Needed as SW scales & complexity increases
- Scale-out engineering application running in cloud





Software Testing Benchmark





The diagram illustrates the architecture of the Software Testing Benchmark. On the left, a 'User' is connected to a 'Cloud master' node, which is represented by a server icon with a cloud above it. The 'Cloud master' is connected to three 'Worker' nodes on the right, each represented by a server icon with a database cylinder next to it. Blue arrows indicate the flow of data and control between the User, the Cloud master, and the three Workers.

- Cloud9, SW Testing as a Service
- Master coordinates symbolic execution
- State maintained in slave, updated from master
- Master load-balances across slaves





CloudSuite 1.0

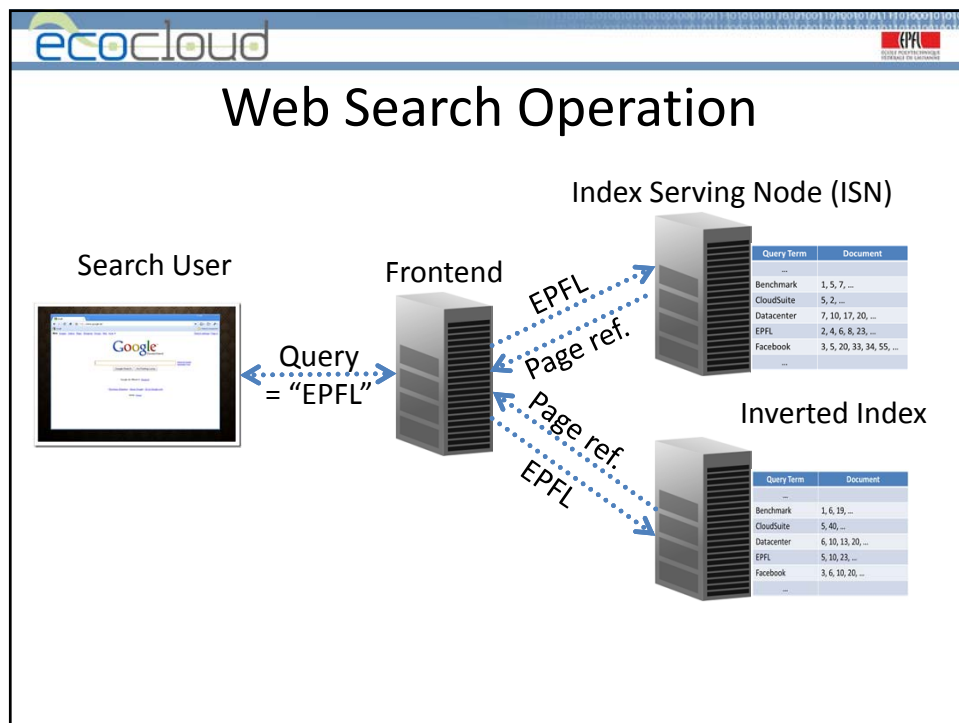
- Data Analytics
- Data Serving
- Media Streaming
- SW Testing
- Web Search
- Web Serving

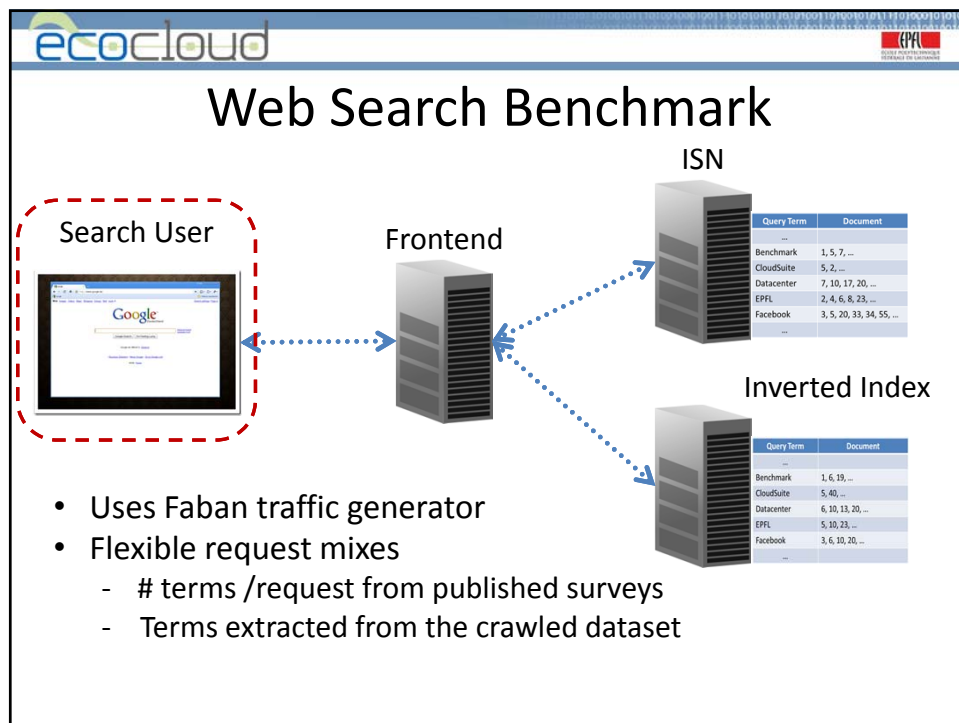
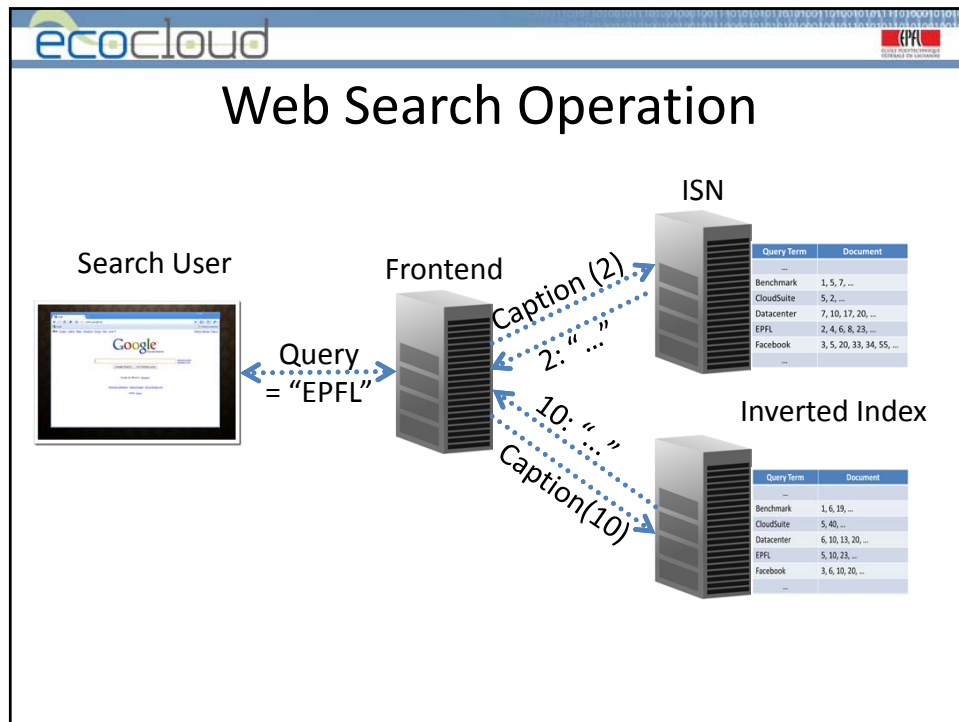


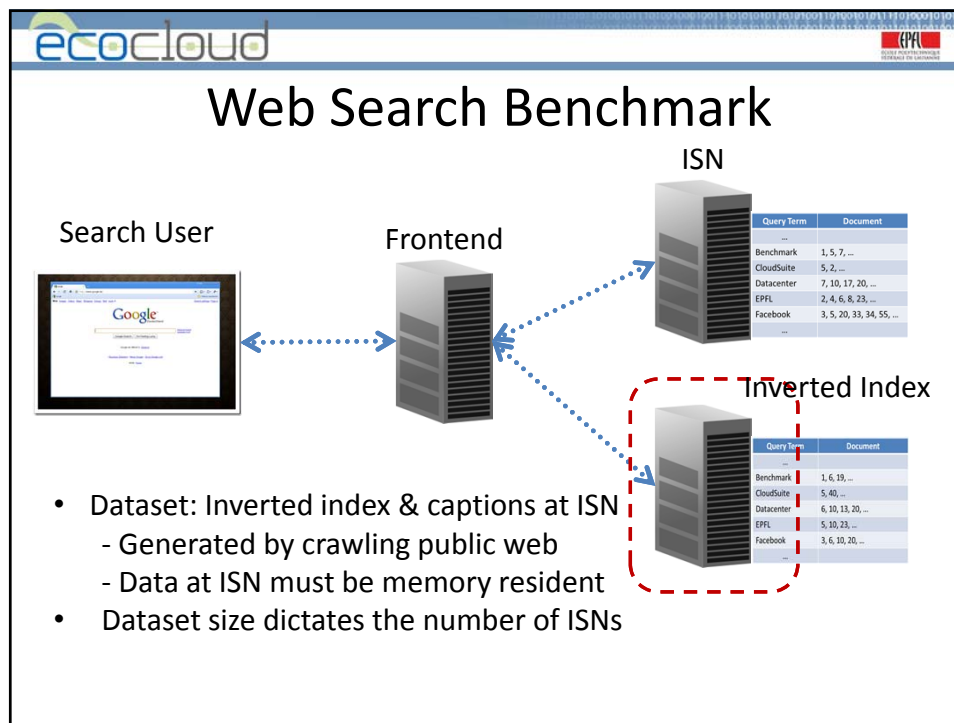
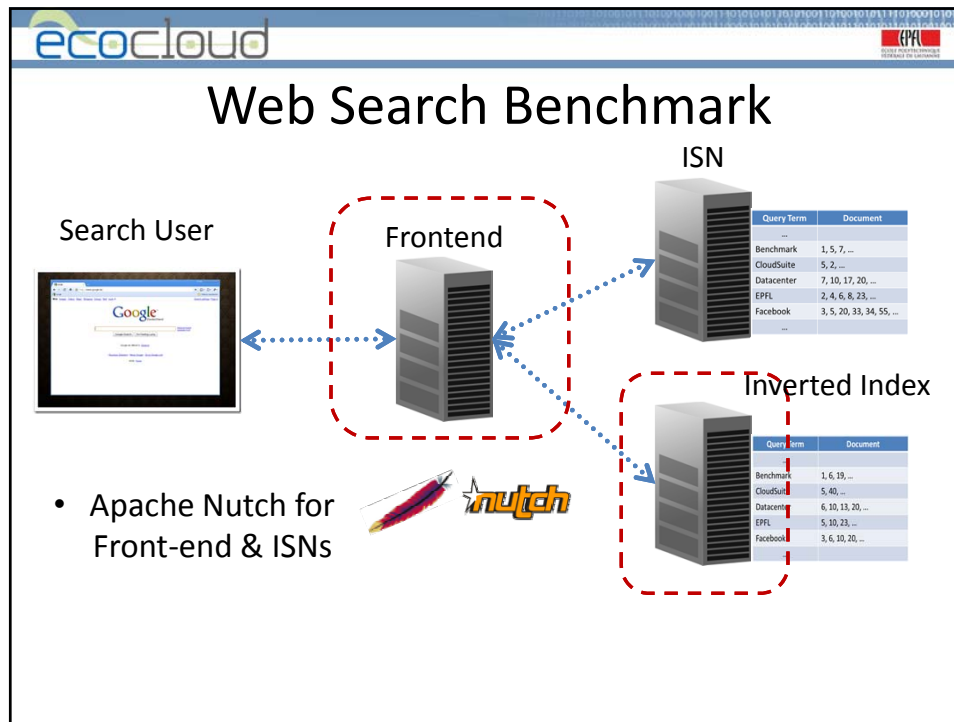
Web Search



- Most popular online service
 - Numerous search clouds deployed by industry












CloudSuite 1.0

- Data Analytics
- Data Serving
- Media Streaming
- SW Testing
- Web Search
- Web Serving




Web Serving

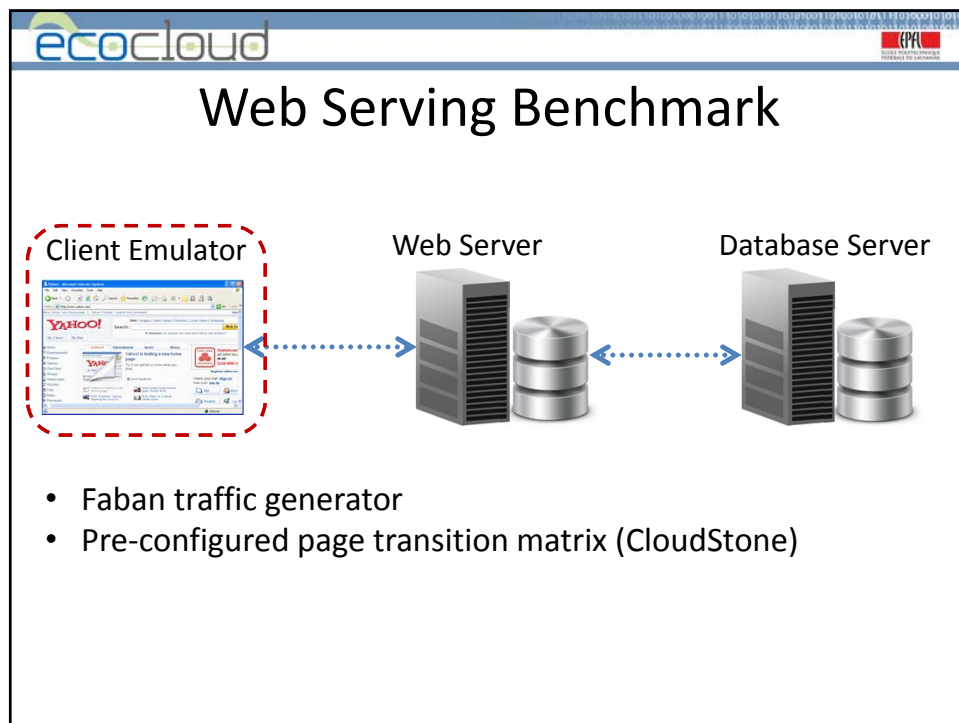
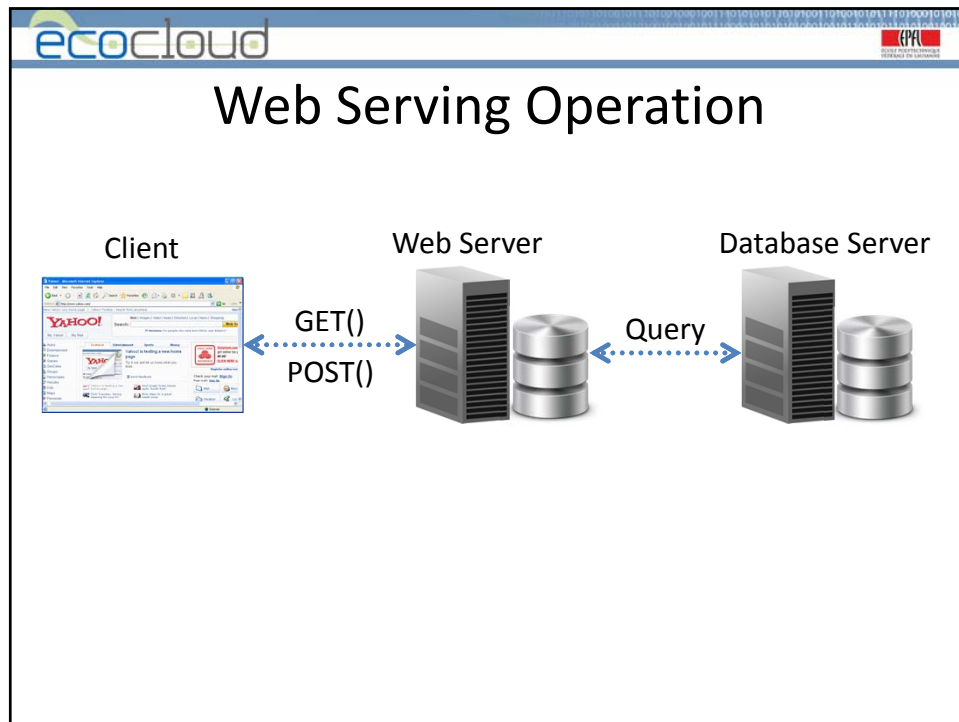
- Key to all internet-based services

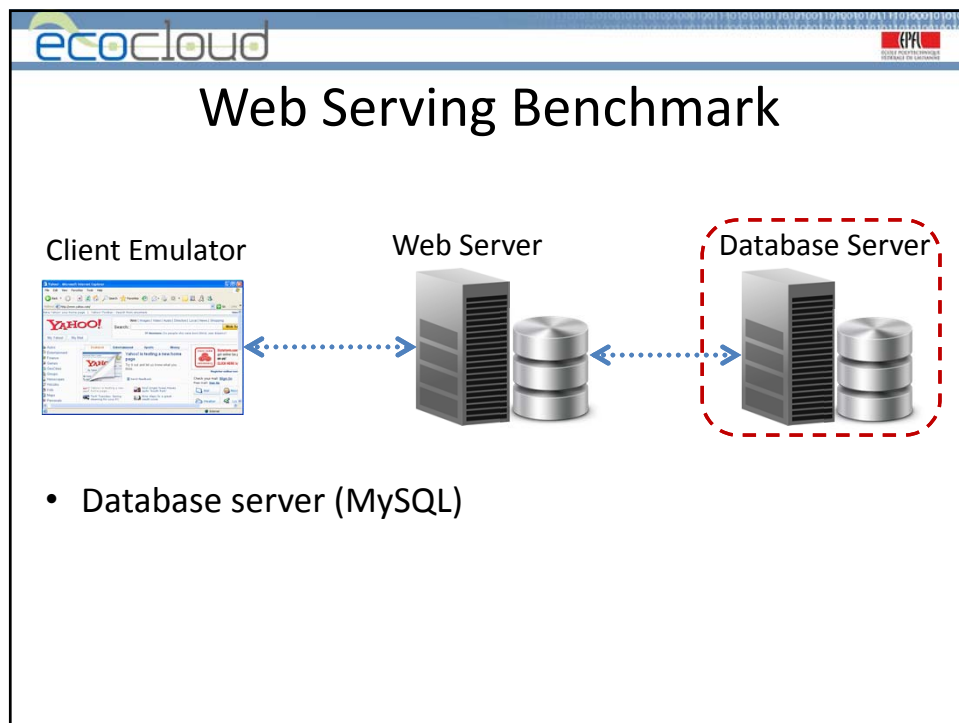
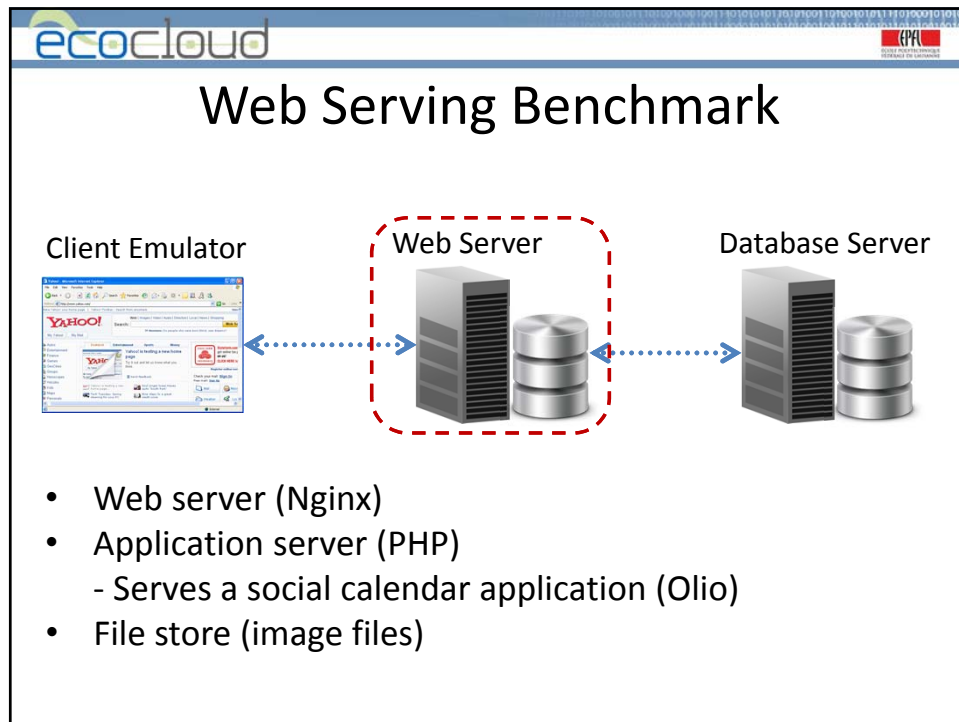




- All services are accessed through web servers



- Various technologies construct web applications
 - HTML, PHP, JavaScript, Ruby









CloudSuite: Hands-on

- Media Streaming
 - Installing the server
 - Installing client generator
 - Overview of the dataset
 - Running the benchmark
 - Checking quality of service



39



CloudSuite Full-System Simulation

Cansu Kaynak

40



CloudSuite Simulation Requirements



CloudSuite Workloads:

- Multi-threaded, multi-processor
- Data-intensive
- Multi-tier

⇒ Exercise OS and I/O extensively
⇒ OS and I/O are first-order performance determinants

Need full-system simulation



41



Flexus Framework

- Functional Full-System Simulation: Simics
- Detailed Microarchitectural Simulation: Flexus
- Fast Simulation: Statistical sampling



42



Flexus Framework

- Functional Full-System Simulation: Simics
- Detailed Microarchitectural Simulation: Flexus
- Fast Simulation: Statistical sampling

43





Full-System Simulation Requirements

Full-system functional simulator must support:

- Privileged-mode ISA
- I/O devices
- Networks of systems
- Saving/restoring architecturally-visible state

Simics provides these capabilities



44



Simics Configuration & CLI

- Configuration file defines system components
 - Motherboard, CPUs, memory, I/O devices
- Command-line interface (CLI) provides interface to simulation
 - Start and stop simulation
 - Save and restore target system checkpoints



45



Simics Checkpoints

- Contain full-system architectural state
- Are incremental - Require all files in chain
- Form the basis for Flexus simulation



46



Simics μ Arch Interface



- Simics does not provide timing details
 - But provides a Micro-Architectural Interface (MAI)
 - Allows a user module to take control over timing
- Simics feeds Flexus with instructions
- Flexus gives timing feedback to Simics

47



Simics Hands-On



48



Preparing a Workload for Simulation

1. Install OS
2. Booting target machine
3. Install application & create dataset
4. Tune workload parameters
5. Run application



49



Preparing a Workload for Simulation

1. Install OS
2. Booting target machine
3. Install application & create dataset
4. Tune workload parameters
5. Run application



50



Media Streaming in Simics Hands-on

1. Loading a freshly-installed OS checkpoint
2. Preparing target system
3. Running applications in Simics
4. Saving system checkpoints
5. Loading system checkpoints



51



Initial Checkpoint

- Freshly-installed OS: Solaris 10 u9
- Media Streaming binaries & datasets
 - Faban client on Client machine
 - Darwin Streaming Server on Server machine
 - Video dataset on Server machine
- Necessary libraries

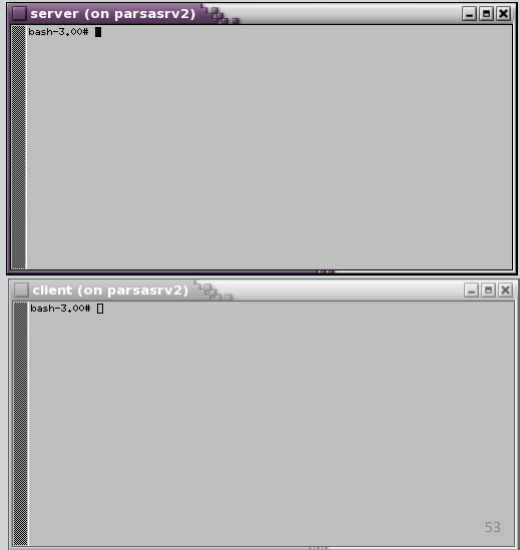
52






Getting Started with Media Streaming

Simulated target system:

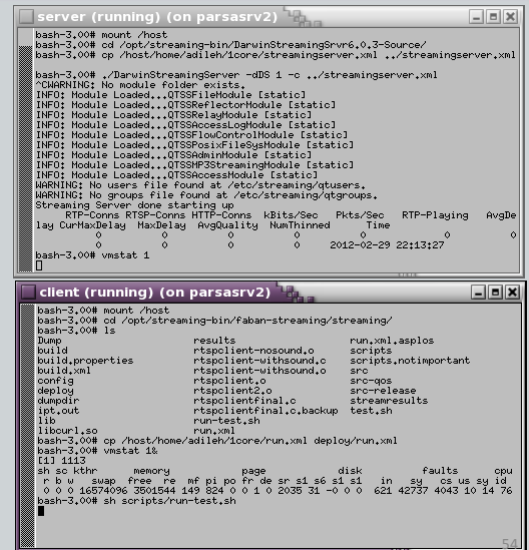
- Server (1 core)
- Client (1 core)
- Binaries:
/opt
- Dataset:
/streaming_data





Preparing Target System

- Move configuration files
- Move experiment files
- Start experiment



Media Streaming in Action

- Monitoring
- QoS check

server (running) (on parsasrv2)

```
INFO: Module Loaded...QTSSAdminModule [static]
INFO: Module Loaded...QTSSMP3StreamingModule [static]
INFO: Module Loaded...QTSSAccessModule [static]
WARNING: No users file found at /etc/streaming/users.
WARNING: No groups file found at /etc/streaming/groups.
Streaming Server done starting up
RTP-Conn RTP-Conn HTTP-Conn
lay CurMaxDelay MaxDelay AvgQuality WBits/Sec Pkts/Sec Time RTP-Playing AvgDe
```

lay	CurMaxDelay	MaxDelay	AvgQuality	WBits/Sec	Pkts/Sec	Time	RTP-Playing	AvgDe
0	0	0	0	0	0	2012-02-29 22:13:48	0	0
0	0	0	0	0	0	2012-02-29 22:13:49	0	0
0	0	0	0	0	0	2012-02-29 22:13:50	0	0
0	0	0	0	0	0	2012-02-29 22:13:51	0	0
0	0	0	0	0	0	2012-02-29 22:13:52	0	0
0	0	0	0	0	0	2012-02-29 22:13:53	0	0
0	0	0	0	0	0	2012-02-29 22:13:54	0	0
0	0	0	0	0	0	2012-02-29 22:13:55	0	0

client (running) (on parsasrv2)

```
1 0 0 70482592 3264192 756 2460 0 0 0 0 0 0 0 0 0 0 386 8061 751 29 5 66
0 0 0 70431688 3247816 755 2441 0 0 0 0 0 0 0 0 0 0 424 8074 753 28 3 69
Feb 29, 2012 10:13:48 PH com.sun.faban.driver.engine.MasterImpl runBenchmark
INFO: RunID for this run is : 1
Feb 29, 2012 10:13:48 PH com.sun.faban.driver.engine.MasterImpl runBenchmark
INFO: Output directory for this run is : /opt/streaming-bin/faban-streaming/tmp/
output/1
Feb 29, 2012 10:13:48 PH com.sun.faban.driver.engine.MasterImpl configureAgents
INFO: Configuring 10 StreamDriverAgents...
0 0 0 70380784 3231480 754 2043 0 0 0 0 0 0 0 0 0 0 457 10672 1259 59 7 34
3 0 0 70133994 3191144 66 1194 0 0 0 0 0 0 0 0 0 0 400 14214 1701 96 2 2
19 0 0 70130640 3181664 13 337 0 0 0 0 0 0 0 0 0 0 524 3442 1324 23 1 76
Kbhr memory page disk faults cpu
r b w swap free re nf pi po fr de sr sl s6 sl s1 in sy cs us sy id
0 0 0 70129136 3179752 0 1 0 0 0 0 0 0 26 0 0 0 658 2054 1176 1 0 99
1 0 0 70129128 3179744 10 592 0 0 0 0 0 0 0 0 0 0 532 1944 589 31 1 68
15 0 0 70126952 3174424 11 748 0 0 0 0 0 0 0 0 0 0 432 13588 556 99 1 0
9 0 0 70122184 3166008 122 9704 0 0 0 0 0 0 0 0 0 0 567 68683 3360 51 19 0
10 0 0 70065128 3094344 166 13091 0 0 0 0 0 0 0 0 0 0 497 47630 6238 42 58 0
13 0 0 70024552 3042152 306 15158 0 0 0 0 0 0 0 0 0 0 443 53032 9987 27 73 0
10 0 0 69947376 3007976 361 16758 0 0 0 0 0 0 0 0 0 0 426 49148 12339 20 80 0
10 0 0 69941976 2974720 368 16308 0 0 0 0 0 0 0 0 0 0 418 45512 12647 22 78 0
13 0 0 69961216 2949960 375 17314 0 0 0 0 0 0 0 0 0 0 428 54942 13106 24 76 0
```






Flexus Simulator Toolset

Stavros Volos





ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE



Software Simulation

- Allows for fast & easy evaluation of an idea
 - Minimal cost, simulator runs on your desktop
 - Reuse components, don't implement everything
- Enables various benchmarks (e.g., SPEC, CloudSuite)
 - Can execute real applications
 - Can simulate thousands of disks
 - Can simulate very fast networks



57



Main Idea

- Use existing system simulator (Simics)
 - Handles BIOS (booting, I/O, interrupt routing, etc...)
- Build a “plugin” architectural model simulator
 - Fast – read state of system from Simics
 - Detailed – interact with and throttle Simics



58



Developing with Flexus

- Flexus philosophy
- Fundamental abstractions
- Important support libraries
- Simulators and components in Flexus 4.1
- Hands-on


59



Flexus philosophy

- Component-based design
 - Compose simulators from encapsulated components
- Software-centric framework
 - Flexus abstractions are not tied to hardware
- Cycle-driven execution model
 - Components receive “clock-tick” signal every cycle
- SimFlex methodology
 - Designed-in fast-forwarding, checkpointing, statistics

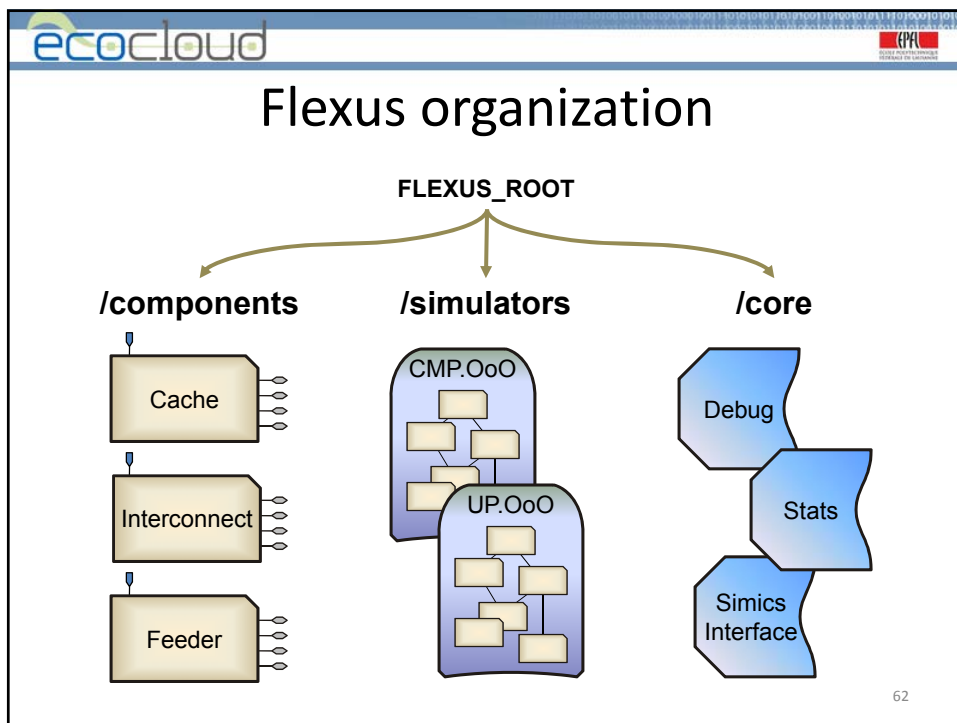
60





Developing with Flexus

- Flexus philosophy
- **Fundamental abstractions**
- Important support libraries
- Simulators and components in Flexus 4.1
- Hands-on

61







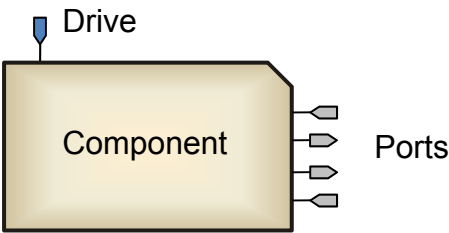
Fundamental abstractions

- Component
 - Component interface
 - Specifies data and control entry points
 - Component parameters
 - Configuration settings available in Simics or cfg file
- Simulator
 - Wiring
 - Specifies which components and how to connect
 - Specifies default component parameter settings

63





Component interface



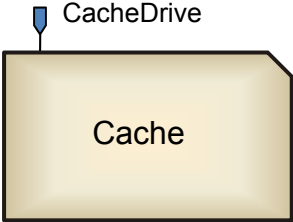
- Component interface (terminology inspired by *Asim* [Emer 02])
 - Drive: “clock-tick” control entry point to component
 - Port: specifies data flow between components

Components w/ same ports are interchangeable

64



Abstractions: Drive





```
classDiagram
    class Cache {
        CacheDrive
    }
```

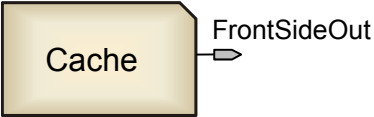
```
COMPONENT_INTERFACE(
    ...
    DRIVE ( Name )
    ...
);
```

- Control entry-point
- Function called once per cycle

65



Abstractions: Port





```
classDiagram
    class Cache {
        FrontSideOut
    }
```

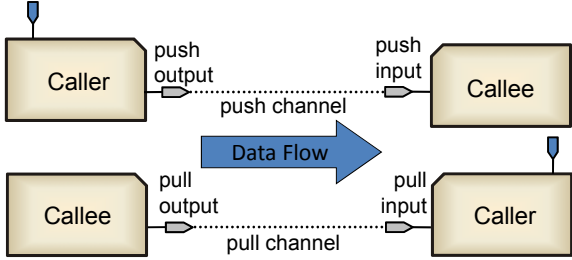
```
COMPONENT_INTERFACE(
    ...
    PORT ( Type, Payload, Name )
    ...
);
```

- Data exchange between components
- Ports connected together in simulator wiring

66





Types of ports and channels

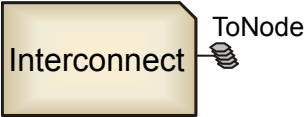


- Type - direction of data and control flow
 - Control flow: Push vs. Pull
 - Data flow: Input vs. Output
- Payload - arbitrary C++ data type
- Type and payload must match to connect ports
- Availability - caller must check if callee is ready

67





Port and component arrays



```
COMPONENT_INTERFACE(  
    ...  
    DYNAMIC_PORT_ARRAY(...)   
    ...  
);
```

- 1-to- n and n -to- n connections
 - E.g., 1 interconnect -> n network interfaces
- Array dimensions can be dynamic

68



Example code using a port



SenderComponent.cpp

```
void someFunction() {  
    Message msg;  
    if ( FLEXUS_CHANNEL(Out).available() ) {  
        FLEXUS_CHANNEL(Out) << msg;  
    }  
}
```

ReceiverComponent.cpp

```
bool available( interface::In )    { return true; }  
void push( interface::In, Message & msg) { ... }
```



69



Configuring components

- Configurable settings associated with component
 - Declared in component specification
 - Can be std::string, int, long, long long, float, double, enum
 - Declaration: `PARAMETER(BlockSize, int, "Cache block size", "bsize", 64)`
 - Use: `cfg.BlockSize`
- Each component instance associated with **configuration**
 - Configuration declared, initialized in simulator wiring file
 - Complete name is `<configuration name>:<short name>`
- Usage from Simics console
 - `flexus.print-configuration` `flexus.write-configuration "file"`
 - `flexus.set "-L2:bsize" "64"`

70



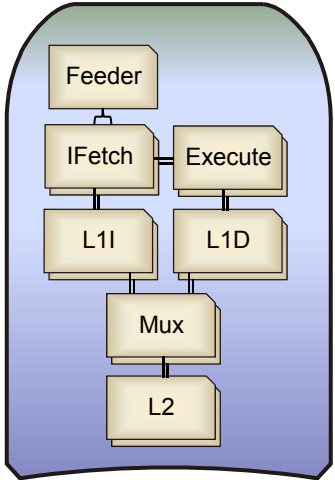
Simulator wiring

[simulators/name/Makefile.name](#)

- List components for link
- Indicate target support



[simulators/name/wiring.cpp](#)

1. Include interfaces
2. Declare configurations
3. Instantiate components
4. Wire ports together
5. List order of drives



```
graph TD; Feeder --> IFetch; IFetch --> L1I; L1I --> Mux; Execute --> L1D; L1D --> Mux; Mux --> L2
```



71



Developing with Flexus

- Flexus philosophy
- Fundamental abstractions
- **Important support libraries**
- Simulators and components in Flexus 4.1
- Hands-on



72



Critical support libraries in /core

- Statistics support library
 - Record results for use with `stat-manager`
- Debug library
 - Control and view Flexus debug messages

73





Statistics support library

- Implements all the statistics you need
 - Histograms
 - Unique counters
 - Instance counters
 - etc...
- Example:

```
Stat::StatCounter myCounter( statName() + "-count" );  
++ myCounter;
```



74



A typical debug statement

```
DBG_(Iface,                                     Severity level
      Comp(*this)                               Associate with this component
      AddCategory( Cache )                     Put this in the "Cache" category
      ( << "Received on FrontSideIn[0](Request): "
        << *(aMessage[MemoryMessageTag])
      )                                         Text of the debug message
      Addr(aMessage[MemoryMessageTag]->address())
    );                                         Add an address field for filtering
```



75



Debug severity levels

1. Tmp	temporary messages (cause warning)
2. Crit	critical errors
3. Dev	infrequent messages, e.g., progress
4. Trace	component defined – typically tracing
5. Iface	all inputs and outputs of a component
6. Verb	verbose output from OoO core
7. Vverb	very verbose output of internals



76



Controlling debug output

- Compile time
 - `make target-severity`
 - (e.g. `make UP.Trace-iface`)
- Run time
 - `flexus.debug-set-severity severity`
- Hint – when you need a lot of detail...
 - Set severity low
 - Run until shortly before point of interest (or failure)
 - Set severity high
 - Continue running



77



Developing with Flexus

- Flexus philosophy
- Fundamental abstractions
- Important support libraries
- **Simulators and components in Flexus 4.1**
- Hands-on



78



Simulators in Flexus 4.1

• UP.Trace	fast memory system
• CMP.L2Shared.Trace	fast CMP memory system
• CMP.MT4.L2Shared.Trace	fast CMP memory system w/ 4-way MT support
• UP.OoO	1 CPU 2-level hierarchy
• CMP.L2SharedNUCA.OoO	private L1 / shared L2
• CMP.MT4.L2SharedNUCA.OoO	private L1 / shared L2 w/ 4-way MT support
• CMP.L2SharedNUCA.DRAMSim.OoO	private L1 / shared L2 w/ DRAMSim 2.0



79



Memory hierarchy

- “top”, “front” = closer to CPU
- Allows for high MLP
 - Non-blocking, pipelined accesses
 - Hit-under-miss within set
- Coherence protocol support
 - Valid, modifiable, dirty states
 - Explicit “dirty” token tracks newest value
 - Non-inclusive
 - Supports “Downgrade” and “Invalidate” messages
 - Request and snoop virtual channels for progress guarantees



80



Out-of-order execution

- Timing-first simulation approach [Mauer'02]
 - OoO components interpret SPARC ISA
 - Flexus validates its results with Simics
- Idealized OoO to maximize memory pressure
 - Decoupled front-end
 - Precise squash & re-execution
 - Configurable ROB, LSQ capacity; dispatch, retire rates
- Memory consistency models (SC, TSO, RMO)

81



Hands-on

- Set up `.run_job.rc.tcl` file
- Launch Simics using the `run_job` script
- Build Flexus simulators
 - Examine Flexus directory structure and source files
- Launch trace-based simulation
- Launch cycle-accurate (OoO) simulation
 - Examine debug output and statistics

How fast is cycle-accurate timing simulation?

82



Flexus: The Way to Full-System Cycle-Accurate Simulation					
Simulator	Full-system Simulation Support	Simulation Speeds			Statistical Sampling Support
		Fast (MIPS)	Trace (KIPS)	Timing (KIPS)	
Flexus	✓	30-60	750-850	20-25	✓
gem5	✓	3-5	35-350	N/A	×
ISA-only	×	Do not support full-system simulation			

ecocloud

CPA


Boosting Simulation Speed with Statistical Sampling

Pejman Lotfi-Kamran





Simulation Speed Challenges

- Longer benchmarks
 - SPEC 2006: Trillions of instructions per benchmark
- Slower simulators
 - Full-system simulation: 1000× slower than SimpleScalar
- Multiprocessor systems
 - CMP: 2x cores every processor generation



1,000,000× slowdown vs. HW → years per experiment

85





Full-system simulation is slow

- Simulation slowdown per cpu

– Real HW:	~ 2 GIPS	1 s
– Simics:	~ 30 MIPS	66 s
– Flexus, no timing:	~ 900 KIPS	37 m
– Flexus, OoO:	~ 24 KIPS	23 h

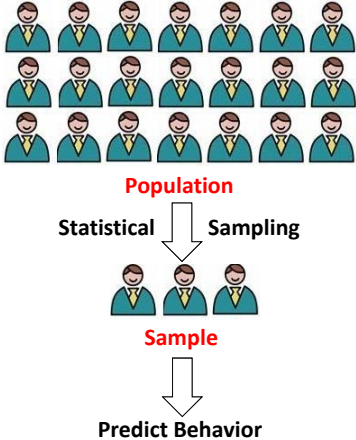
2 years to simulate 10 seconds of a 64-core workload!

86



Statistical Sampling

- Random selection of population
 - E.g., 3000 out of 300 million
- Predict the behavior based on the selected sample
- Features:
 - High accuracy
 - Simple
 - Strong mathematical foundation



Population



Statistical Sampling

Sample

Predict Behavior


Power of a small part to predict behavior of a whole

87



Statistical Sampling for Simulation

- Measure uniform or random locations





measurements

- Each measurement is on a group of instructions
- ~10,000x reduction in turnaround time

Challenge: programs are sequential



88



Sampling of Sequential Programs


- Correctness
 - State of memory, registers, etc.
- Bias
 - State of cache, branch predictor, reorder buffer, etc.

89



Functional Simulation



- Functional simulation is faster than detailed simulation
 - Flexus (no timing) is 38 times faster than Flexus (OoO)
- Use functional simulation for “warmup”
 - Memory (guarantees correctness)
 - Cache hierarchy (avoids bias)
 - Branch predictor (avoids bias)



— Functional warming ■ Measurement


No state for core microarchitecture → Bias

90





Handling Bias

- Core micro-architecture can be warmed up rapidly
 - Detailed simulation to warmup core micro-architecture
- Perform warmup prior to measurement
 - Functional warming during fast-forwarding
 - Detailed warmup before each simulation window

SMARTS 


— Functional warming — Detailed warmup — Measurement

91





Simulation Speedup

- 10 seconds of a 64-core workload
 - Normal execution: 2 years
 - With sampling: 20 days

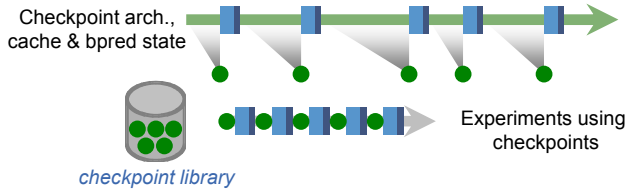


- 37x improvement in simulation speed but not enough
- Solution
 - Avoid functional simulation (17 days)
 - Accelerate detailed simulation (3 days)

92





Avoiding Functional Simulation



- Store warm cache & branch predictor state
 - Same sample design, accuracy, confidence
 - No warming length prediction needed


Works for any microarchitecture

93





Accelerating Detailed Simulation

- Checkpoint library makes measurement independent
- Run multiple measurements in parallel




Run in parallel

94




Simulation Speedup



- Sampling without a checkpoint library:
 - 10 seconds of a 64-core workload: 20 days



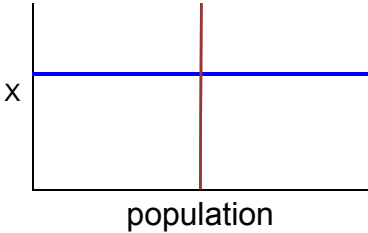
- Sampling with a checkpoint library:
 - 10 seconds of a 64-core workload: 3 hours with 100 cores



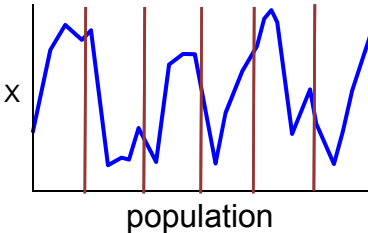
95



How to Choose the Sample Size?





Low variability → Small sample size






High variability → Large sample size

Variability determines sample size



96



Steps for Timing Simulation

1. Prepare workload for simulation
 - Port workload into Simics
2. Measure baseline variance
 - Determine required library size
3. Collect checkpoints
 - Via functional warmup
4. Detailed Simulation
 - Estimate performance results

97



2. Determine Sampling Parameters

- Guess variability
- Generate flexpoints for the variability
- Run timing simulation
- Measure error and correct the guess

98

Typical Sampling Parameters	
	Flexus (64-CPU CMP.OoO)
Warming	100k cycles
Measurement	50k cycles
Target confidence	95%
Sample size	800
Sim. time per checkpoint	~ 20 min

99



3. Checkpoint Creation

- Spread Simics checkpoints
 - Simics fast mode rapidly covers 10 seconds
- Collect flexpoints in parallel
 - Via CMP.L2Shared.Trace
 - From each Simics checkpoint

Simics checkpoint, "Phase"

Simics + Flexus checkpoint, "Flexpoint"



100



4. Detailed Simulation

- Run detailed simulation with OoO simulators
- Process all flexpoints, aggregate offline
- Manipulate results with *stat-manager*
 - Each run creates binary `stats_db.out` database
 - Offline tools to select subsets; aggregate
 - Generate text reports from simple templates
 - Compute confidence intervals for mean estimates

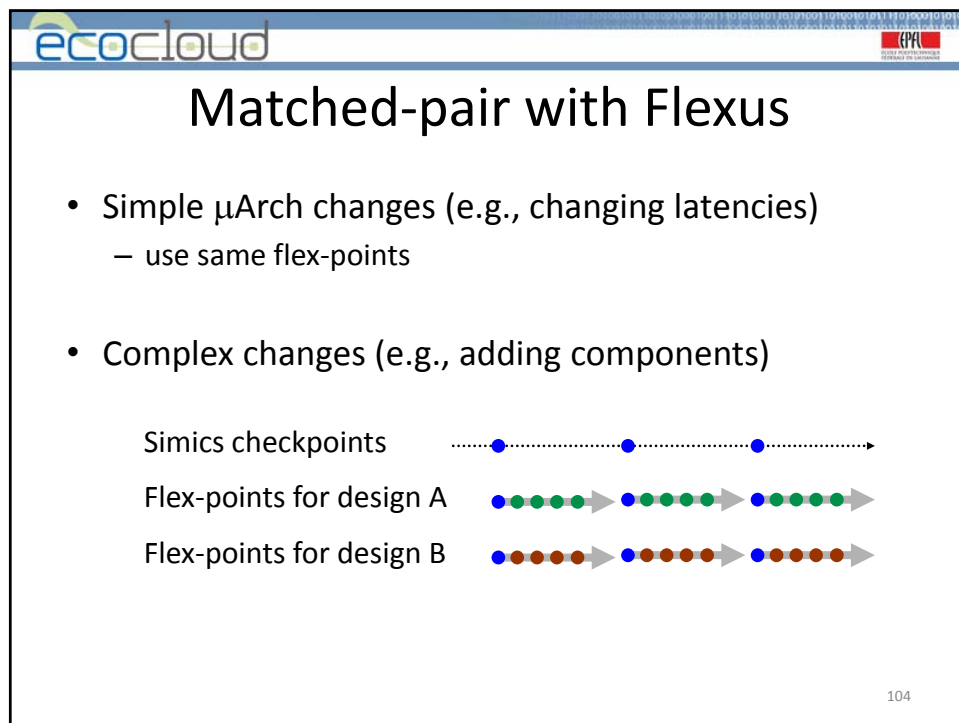
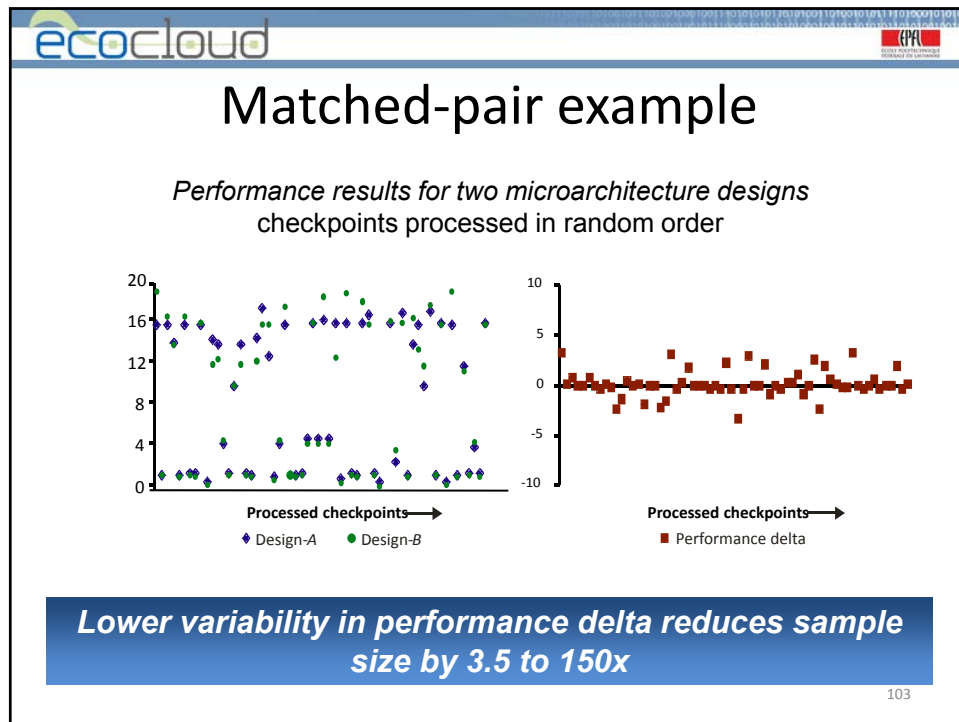
101





Matched-pair comparison [Ekman 05]

- Often interested in relative performance
- Change in performance across designs varies less than absolute change
- Matched pair comparison
 - Allows smaller sample size
 - Reports confidence in performance change

102







Hands-on

- Generate Flexpoints
- Launch timing simulation for all flexpoints
- Aggregate stats with stat-collapse
- Examine aggregate statistics
 - Compute confidence
 - Plot timing breakdown



105



How to Use CloudSuite Images

Cansu Kaynak

106





CloudSuite Simics Release

Released images (phase_000) contain:

- CloudSuite binaries & necessary libraries
- Tuned workloads at steady state
- Ready to run

107





CloudSuite Images

From 1 core to 64 cores:

1. Data Analytics
2. Data Serving
3. Media Streaming (4, 8, 16 cores)
4. Software Testing
5. Web Search (1 to 32 cores) ~ SW scalability
6. Web Serving (1 to 8 cores)

108



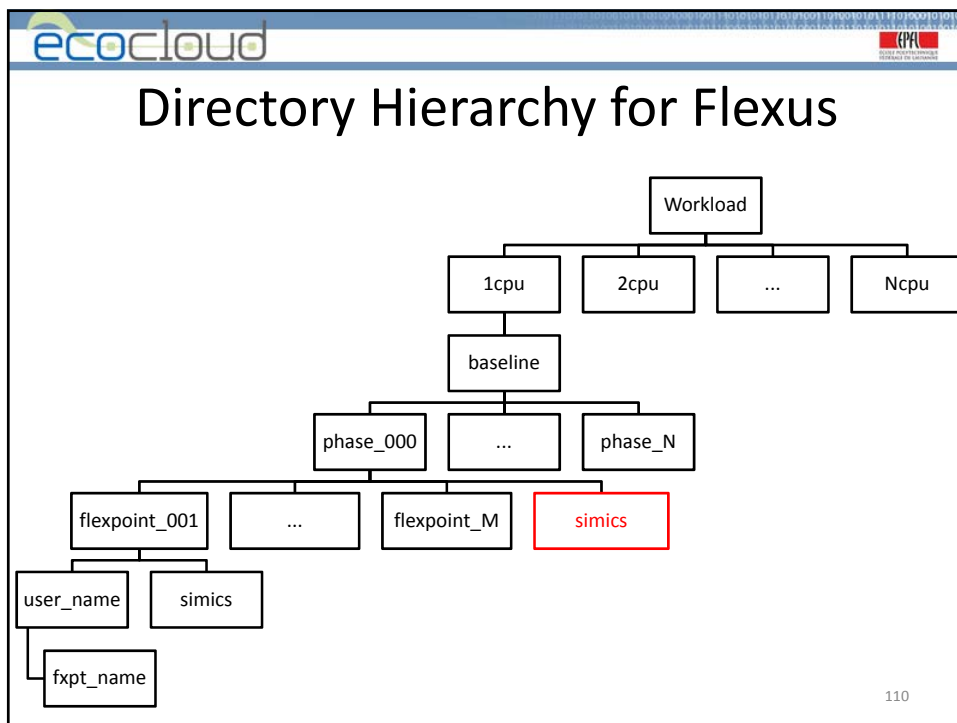
Deploying CloudSuite Images



- Paths for logical components in configuration files:
 - Binary disk
 - Data disk(s)

```
checkpoint_path: (  "/path/to/binary_disk",  
                  "/path/to/data_disk")
```

- Load initial state & save it as phase_000
- Detailed instruction are in setup document...

109








What We Release

We provide phase_000:

- Steady state of workload execution

Execution ... 

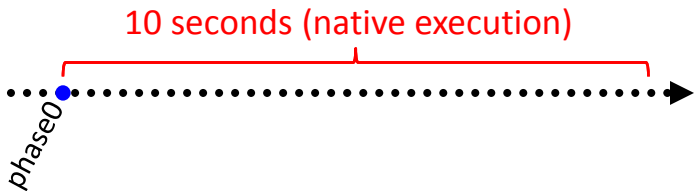
111





How Long To Simulate

Representative execution window of a workload:

- Steady architectural behavior (measured on real HW)
- 10 sec. of native execution (25 sec. for media streaming)

Execution ... 

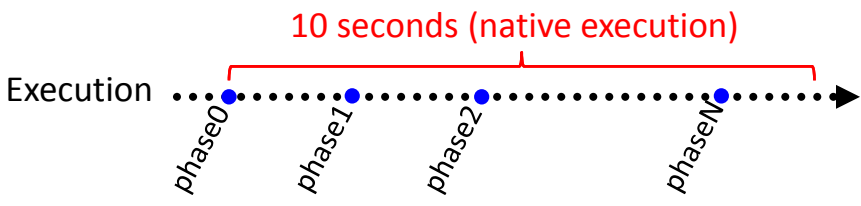
112



Phase Generation



Divides the entire execution into phases

- Generates phases (Simics checkpoints) using Simics fast mode
- As many phases as necessary for desired parallelism
 - e.g., 10 phases



The diagram shows a horizontal timeline labeled 'Execution' on the left and an arrow pointing right on the right. A series of dots represents the execution flow. Four blue dots are marked and labeled 'phase0', 'phase1', 'phase2', and 'phaseN' from left to right. A red bracket spans the distance between 'phase0' and 'phaseN', with the text '10 seconds (native execution)' written above it in red.

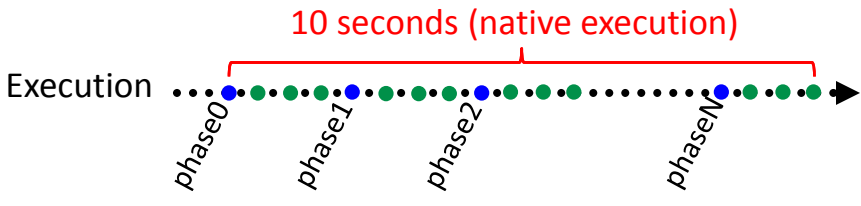
113



Flexpoint Generation



Divides every phase into flexpoints (parallel across phases)

- Generates flexpoints using Flexus trace simulator
 - Functional warming of cache and branch predictor state
- As many flexpoints as necessary for desired degree of confidence
 - e.g., 80 flexpoints per phase



The diagram shows a horizontal timeline labeled 'Execution' on the left and an arrow pointing right on the right. A series of dots represents the execution flow. Four blue dots are marked and labeled 'phase0', 'phase1', 'phase2', and 'phaseN' from left to right. Between each blue dot, there are several green dots, representing flexpoints. A red bracket spans the distance between 'phase0' and 'phaseN', with the text '10 seconds (native execution)' written above it in red.

114




Timing Simulation



Cycle-accurate simulation in parallel across flexpoints

- First, detailed warm-up of microarchitectural state
- Then, takes measurements from the warmed state
 - e.g., 100K-cycle warm-up, 50K-cycle measurement
 - Longer warm-up necessary for Data Serving

Independent parallel simulations

Execution ... 

115



Wrap-Up

- Two steps before cycle-accurate simulation:
 1. Phase generation
 2. Flexpoint generation
- Refer to `.run_job.rc.tcl` in Flexus 4.1 for workloads, phases, flex-points

116

